

Home > Cheat sheets > Machine Learning

Scikit-Learn Cheat Sheet: Python Machine Learning

A handy scikit-learn cheat sheet to machine learning with Python, including some code examples.

May 2021 · 4 min read



Karlijn Willems

Former Data Journalist at DataCamp | Manager at NextWave Consulting

TOPICS

Machine Learning

Python

Most of you who are learning data science with Python will have definitely heard already about scikit-learn, the open source Python library that implements a wide variety of machine learning, preprocessing, cross-validation and visualization algorithms with the help of a unified interface.

If you're still quite new to the field, you should be aware that machine learning, and thus also this Python library, belong to the must-knows for every aspiring data scientist.

That's why DataCamp has created a scikit-learn cheat sheet for those of you who have already started learning about the Python package, but that still want a handy reference sheet. Or, if you still have no idea about how scikit-learn works, this machine learning cheat sheet might come in handy to get a quick first idea of the basics that you need to know to get started.

Either way, we're sure that you're going to find it useful when you're tackling machine learning problems!

This scikit-learn cheat sheet will introduce you to the basic steps that you need to go through to implement machine learning algorithms successfully: you'll see how to load in your data, how to preprocess it, how to create your own model to which you can fit your data and predict target labels, how to validate your model and how to tune it further to improve its performance.



Have this cheat sheet at your fingertips

[↓ Download PDF](#)

In short, this cheat sheet will kickstart your data science projects: with the help of code examples, you'll have created, validated and tuned your machine learning models in no time.

So what are you waiting for? Time to get started!

(Click above to download a printable version or read the online version below.)

Python For Data Science Cheat Sheet: Scikit-learn

Scikit-learn is an open source Python library that implements a range of machine learning, preprocessing, cross-validation and visualization algorithms using a unified interface.

A Basic Example

```
from sklearn import neighbors, datasets, preprocessing
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
iris = datasets.load_iris()
X, y = iris.data[:, :2], iris.target
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=33)
scaler = preprocessing.StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
knn = neighbors.KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
accuracy_score(y_test, y_pred)
```

POWERED BY DATACAMP WORKSPACE

Loading The Data

Your data needs to be numeric and stored as NumPy arrays or SciPy sparse matrices. Other types that are convertible to numeric arrays, such as Pandas DataFrame, are also acceptable.

```
import numpy as np
X = np.random.random((10,5))
y = np.array(['M', 'M', 'F', 'F', 'M', 'F', 'M', 'M', 'F', 'F'])
X[X < 0.7] = 0
```

POWERED BY DATACAMP WORKSPACE

Preprocessing The Data

Standardization

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler().fit(X_train)
standardized_X = scaler.transform(X_train)
standardized_X_test = scaler.transform(X_test)
```

POWERED BY DATACAMP WORKSPACE

Normalization

```
from sklearn.preprocessing import Normalizer
scaler = Normalizer().fit(X_train)
normalized_X = scaler.transform(X_train)
normalized_X_test = scaler.transform(X_test)
```

POWERED BY DATACAMP WORKSPACE

Binarization

```
from sklearn.preprocessing import Binarizer
binarizer = Binarizer(threshold=0.0).fit(X)
binary_X = binarizer.transform(X)
```

POWERED BY DATACAMP WORKSPACE

Encoding Categorical Features

```
from sklearn.preprocessing import LabelEncoder
enc = LabelEncoder()
y = enc.fit_transform(y)
```

POWERED BY DATACAMP WORKSPACE

Imputing Missing Values

```
from sklearn.preprocessing import Imputer
imp = Imputer(missing_values=0, strategy='mean', axis=0)
imp.fit_transform(X_train)
```

POWERED BY DATACAMP WORKSPACE

Generating Polynomial Features

```
from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(5)
poly.fit_transform(X)
```

POWERED BY DATACAMP WORKSPACE

Training And Test Data

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

POWERED BY DATACAMP WORKSPACE

Create Your Model

—

**Final chance: up to 67% off
annual subscriptions**
Offer ends in 1d 21h 25m 22s



```
from sklearn.linear_model import LinearRegression
lr = LinearRegression(normalize=True)
```

POWERED BY DATACAMP WORKSPACE

Support Vector Machines (SVM)

```
from sklearn.svm import SVC
svc = SVC(kernel='linear')
```

POWERED BY DATACAMP WORKSPACE

Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
```

POWERED BY DATACAMP WORKSPACE

KNN

```
from sklearn import neighbors
knn = neighbors.KNeighborsClassifier(n_neighbors=5)
```

POWERED BY DATACAMP WORKSPACE

Unsupervised Learning Estimators

Principal Component Analysis (PCA)

```
from sklearn.decomposition import PCA
pca = PCA(n_components=0.95)
```

POWERED BY DATACAMP WORKSPACE

K Means

```
from sklearn.cluster import KMeans
```

```
k_means = KMeans(n_clusters=3, random_state=0)
```

POWERED BY DATACAMP WORKSPACE

Model Fitting

Supervised learning

```
lr.fit(X, y)
knn.fit(X_train, y_train)
svc.fit(X_train, y_train)
```

POWERED BY DATACAMP WORKSPACE

Unsupervised Learning

```
k_means.fit(X_train)
pca_model = pca.fit_transform(X_train)
```

POWERED BY DATACAMP WORKSPACE

Prediction

Supervised Estimators

```
y_pred = svc.predict(np.random.random((2,5)))
y_pred = lr.predict(X_test)
y_pred = knn.predict_proba(X_test)
```

POWERED BY DATACAMP WORKSPACE

Unsupervised Estimators

```
y_pred = k_means.predict(X_test)
```

POWERED BY DATACAMP WORKSPACE

Evaluate Your Model's Performance

Classification Metrics

Accuracy Score

```
knn.score(X_test, y_test)
from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred)
```

POWERED BY DATACAMP WORKSPACE

Classification Report

```
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

POWERED BY DATACAMP WORKSPACE

Confusion Matrix

```
from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred))
```

POWERED BY DATACAMP WORKSPACE

Regression Metrics

Mean Absolute Error

```
from sklearn.metrics import mean_absolute_error
y_true = [3, -0.5, 2]
mean_absolute_error(y_true, y_pred)
```

POWERED BY DATACAMP WORKSPACE

Mean Squared Error

```
from sklearn.metrics import mean_squared_error
mean_squared_error(y_test, y_pred)
```

POWERED BY DATACAMP WORKSPACE

R2 Score

```
from sklearn.metrics import r2_score
r2_score(y_true, y_pred)
```

POWERED BY DATACAMP WORKSPACE

Clustering Metrics

Adjusted Rand Index

```
from sklearn.metrics import adjusted_rand_score
adjusted_rand_score(y_true, y_pred)
```

POWERED BY DATACAMP WORKSPACE

Homogeneity

```
from sklearn.metrics import homogeneity_score
homogeneity_score(y_true, y_pred)
```

POWERED BY DATACAMP WORKSPACE

V-measure

```
from sklearn.metrics import v_measure_score
metrics.v_measure_score(y_true, y_pred))
```

POWERED BY DATACAMP WORKSPACE

Cross-Validation

```
print(cross_val_score(knn, X_train, y_train, cv=4))
print(cross_val_score(lr, X, y, cv=2))
```

POWERED BY DATACAMP WORKSPACE

Tune Your Model

Grid Search

```
from sklearn.grid_search import GridSearchCV
params = {"n_neighbors": np.arange(1,3), "metric": ["euclidean", "cityblock"]}
grid = GridSearchCV(estimator=knn, param_grid=params)
grid.fit(X_train, y_train)
print(grid.best_score_)
print(grid.best_estimator_.n_neighbors)
```

POWERED BY DATACAMP WORKSPACE

Randomized Parameter Optimization

```
from sklearn.grid_search import RandomizedSearchCV
params = {"n_neighbors": range(1,5), "weights": ["uniform", "distance"]}
rsearch = RandomizedSearchCV(estimator=knn,
    param_distributions=params,
    cv=4,
    n_iter=8,
    random_state=5)
rsearch.fit(X_train, y_train)
print(rsearch.best_score_)
```

POWERED BY DATACAMP WORKSPACE

Going Further

Begin with [our scikit-learn tutorial for beginners](#), in which you'll learn in an easy, step-by-step way how to explore handwritten digits data, how to create a model for it, how to fit your data to your model and how to predict target values. In addition, you'll make use of Python's data visualization library matplotlib to visualize your results.

PS. Don't miss our [Bokeh cheat sheet](#), the [Pandas cheat sheet](#) or the [Python cheat sheet for data science](#).

TOPICS

Related



How to Install Python

Richie Cotton



Understanding Data Drift and
Model Drift: Drift Detection in...

Moez Ali



How to Create a Histogram with
Plotly

Kurtis Pykes

See More →

Grow your data skills with DataCamp for Mobile

Make progress on the go with our mobile courses and daily 5-minute coding challenges.



LEARN

Learn Python

Learn R

Learn SQL

Learn Power BI

Learn Tableau

Assessments

Career Tracks

Skill Tracks

Courses

Data Science Roadmap

DATA COURSES

Python Courses

R Courses

SQL Courses

Power BI Courses

Tableau Courses

Spreadsheet Courses

Data Analysis Courses

Data Visualization Courses

Machine Learning Courses

Data Engineering Courses

WORKSPACE

Get Started

Templates

Integrations

Documentation

CERTIFICATION

Certifications

Data Scientist

Data Analyst

Hire Data Professionals

RESOURCES

Resource Center

Upcoming Events

Blog

[Tutorials](#)

[Open Source](#)

[RDokumentation](#)

[Course Editor](#)

[Book a Demo with DataCamp for Business](#)

PLANS

[Pricing](#)

[For Business](#)

[For Classrooms](#)

[Discounts, Promos & Sales](#)

[DataCamp Donates](#)

SUPPORT

[Help Center](#)

[Become an Instructor](#)

[Become an Affiliate](#)

ABOUT

[About Us](#)

[Learner Stories](#)

[Careers](#)

[Press](#)

[Leadership](#)

[Contact Us](#)



[Privacy Policy](#)

[Cookie Notice](#)

[Do Not Sell My Personal Information](#)

[Accessibility](#)

[Security](#)

[Terms of Use](#)

© 2023 DataCamp, Inc. All Rights Reserved.